RouteRec: Behavior-Guided Expert Routing for Sequential Recommendation

Anonymous Author(s)

Abstract

Sequential recommendation is commonly optimized with one shared computation pipeline, even though interaction regimes vary substantially across users, sessions, and time scales. In practice, some prefixes are dominated by stable preference, some by repetition, and others by local exploration. This mismatch between behavioral diversity and shared computation is increasingly visible as backbones become stronger: performance improves, but the model still offers little control over *how* different behavior patterns should trigger different computation.

Mixture-of-experts (MoE) is a natural candidate for conditional computation, yet many sequential MoE variants still route primarily from hidden states. This creates a semantic bottleneck: the same latent signal is asked to encode sequence content and decide routing allocation. As a result, expert assignment is often effective but opaque, making it difficult to interpret specialization, diagnose failure modes, or align routing with observable behavioral factors.

We propose **RouteRec**, a behavior-guided MoE framework that explicitly decouples routing signals from representation signals. RouteRec keeps a strong sequence encoder for hidden-state construction, while routing is driven by compact behavior cues extracted from standard logs (timestamps, optional categories, and interaction statistics). These cues are organized into four semantic families (Tempo, Focus, Memory, Exposure) and consumed through three temporal scopes (macro, mid, micro), so the routing structure matches evidence granularity rather than imposing a single gating bias everywhere.

This design is intentionally lightweight: we avoid heavy side-information modeling in the routing branch and instead emphasize interpretable, transferable summaries that are easy to compute across datasets. We further provide a dedicated preliminary formulation to make the method readable without overloading notation. Experiments show that RouteRec improves next-item recommendation over strong baselines, and diagnostic analyses indicate that improvements are associated with more behavior-consistent expert usage rather than unstructured parameter scaling.

CCS Concepts

• **Information systems** → **Recommender systems**.

Keywords

sequential recommendation, mixture of experts, conditional computation, routing, behavioral features

RecSys '26, Minneapolis, MN, USA
2026.

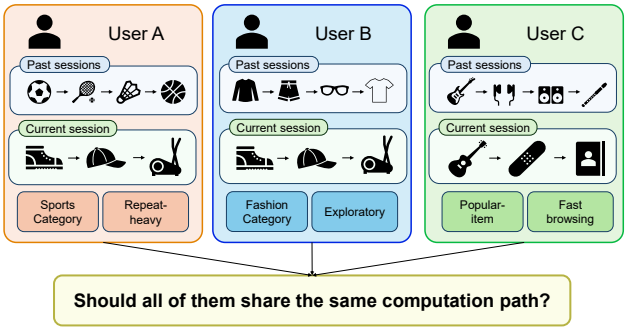

Figure 1: Different users and sessions can require different computational paths.

ACM Reference Format:
Anonymous Author(s). 2026. RouteRec: Behavior-Guided Expert Routing for Sequential Recommendation. In *Proceedings of the 20th ACM Conference on Recommender Systems (RecSys '26), September 28–October 2, 2026, Minneapolis, MN, USA.* ACM, New York, NY, USA, 13 pages.

1 Introduction

Sequential recommendation is usually formulated as next-item prediction over an ordered interaction history. At first glance, this appears to be a homogeneous prediction task, but real user behavior is not homogeneous at all. Even for the same user, one session may reflect persistent preference, another may reflect short-term mission behavior, and another may be dominated by quick exploratory clicks. Prior surveys and studies consistently describe this multi-regime nature of sequential and session-based recommendation [29, 30].

Modern backbones have become very strong at sequence modeling [8, 11, 14, 17, 21, 24], but they still predominantly rely on shared transformations across behaviorally different cases. This design is efficient and often performant on average, yet it leaves an important modeling question open: when two prefixes imply fundamentally different behavioral contexts, should they really invoke the same computation path?

Figure 1 provides an intuitive view of this gap. Distinct users and sessions can exhibit markedly different tendencies in pace, repetition, and local exploration, but a one-path design does not expose a mechanism for behavior-specific computation allocation. In other words, representation quality improves, but routing semantics remain implicit.

Mixture-of-experts (MoE) provides a direct mechanism for conditional computation [5, 9, 10, 23]. However, in many recommendation settings, routing is still derived from hidden states. This is practical, but it also entangles two roles that are

better separated: representing sequence content and deciding where computation should go. When both are handled by the same latent signal, expert specialization becomes harder to interpret and debug.

This limitation becomes more serious when the goal is not only accuracy but also controllability and diagnosis. If routing is purely latent, it is difficult to answer basic questions such as: Which behavior pattern activated which expert? Are similar sessions routed consistently? Is an expert specializing in a meaningful regime or merely in an embedding artifact? These questions are central for robust deployment and scientific understanding, but are weakly supported by hidden-state-only routing.

We therefore adopt a simple principle: **decouple routing signals from representation signals**. In RouteRec, hidden states remain the content channel processed by experts, while routing is controlled by compact behavior cues traceable to logs. This separation preserves backbone strength while making routing decisions easier to inspect and reason about.

RouteRec also aligns routing with temporal scope. Behavior evidence near the prediction point is not equivalent to long-range history evidence; forcing a single routing bias for all horizons is unnecessarily rigid. We therefore use a multi-stage design with macro, mid, and micro routing, allowing coarse-to-fine expert allocation that reflects cross-session tendency, within-session evolution, and immediate transitions.

Our goal is not to maximize notation complexity or introduce heavy side branches. Instead, we emphasize a compact, transferable routing view built from broadly available signals. This makes the approach practical across datasets and improves method readability, while still enabling richer conditional computation than shared-path baselines.

This paper also introduces a dedicated *Preliminary* section before the main method. The goal is pragmatic: define the problem statement, notation, and MoE basics once, so the method section can focus on RouteRec-specific design choices.

The main contributions of this work are as follows:

- We identify a core limitation of hidden-state-only routing in sequential MoE recommendation and formalize routing-representation decoupling as a practical design principle.
- We introduce RouteRec, a behavior-guided multi-stage MoE framework that routes with compact cues from standard interaction logs instead of hidden-state-only gating.
- We provide a clear preliminary formulation (problem statement and MoE background) to support readable method development without excessive symbol overhead.
- We show that RouteRec improves recommendation quality over strong baselines and yields more interpretable, behavior-consistent routing patterns in analysis.

2 Related Work

2.1 Sequential Recommendation Backbones

Sequential recommendation has progressed from recurrent and convolutional models to graph and Transformer architectures [8, 11, 25, 27, 31]. Recent methods further strengthen performance with time-aware attention, contrastive objectives, and structural inductive bias [4, 14, 17, 21, 24, 28, 32]. These advances improve representation quality, but they generally retain shared computation across behaviorally different cases.

2.2 Heterogeneous User State Modeling

Many studies explicitly acknowledge non-stationary behavior, including short-term intent modeling and multi-interest extraction [3, 12, 13, 16, 18, 35]. Most approaches handle heterogeneity inside representation learning, for example by attention reweighting or interest decomposition. RouteRec is complementary: it treats heterogeneity as a routing problem and uses explicit behavior cues to modulate computation paths.

2.3 Feature-Aware Sequential Recommendation

Feature-aware methods use timestamps, side information, and feature channels to strengthen sequence encoding [4, 14, 33, 37]. In most cases, these signals are fused into hidden representations. RouteRec uses a different role assignment: features are primarily routing controls, while representation transformation remains expert-specific over hidden states.

2.4 MoE in Recommendation

Classical and modern MoE work demonstrates the value of conditional computation [5, 9, 10, 23]. Recommendation-oriented expert architectures also show strong capacity benefits [1, 15, 19, 26, 36]. Our distinction is not simply adding experts, but making routing semantics behavior-grounded and interpretable through explicit cue design.

3 Preliminary

3.1 Problem Statement

For each user, we observe interactions ordered by time and train the model for next-item prediction. In one training instance, we denote the observed prefix as $x_{1:t}$ and the target as x_{t+1}. Each interaction includes at least an item identity and timestamp, and may additionally include coarse category information depending on the dataset.

To keep the formulation readable, we use two signal channels throughout this paper. The first is the representation channel, where a sequential encoder maps observed prefixes to hidden states $\mathbf{h}_t$. The second is the behavior channel, where lightweight statistics from the same prefix are summarized into a compact cue vector $\mathbf{b}_t$. RouteRec uses $\mathbf{h}_t$ for expert transformation and $\mathbf{b}_t$ for routing decisions. This separation

is the central modeling choice that the later method section instantiates stage by stage.

The behavior channel is intentionally simple: cues are derived from standard log attributes rather than from a deep side encoder. This keeps routing signals stable across datasets and easier to interpret during analysis. In practice, cues capture aspects such as pace, repetition tendency, focus concentration, and popularity exposure. We do not force a large handcrafted taxonomy in this section; instead, we keep the abstraction compact and defer stage-specific feature construction details to Method.

3.2 MoE Background for Sequential Models

In a standard MoE layer with E experts, an expert-weighted output at position t is

$$\tilde{\mathbf{h}}_t = \sum_{e=1}^{E} \pi_{t,e} \operatorname{Expert}_e(\mathbf{h}_t), \tag{1}$$

where $\boldsymbol{\pi}_t \in \mathbb{R}^E$ is a routing distribution with $\sum_e \pi_{t,e} = 1$.

In many practical systems, $\boldsymbol{\pi}_t$ is computed from hidden states. This is effective but often semantically opaque, because routing and content encoding are entangled in the same latent space. RouteRec follows a different view: routing is computed from behavior cues, while hidden states are dedicated to expert transformation. The distinction is conceptually small but operationally important for interpretability and diagnosis.

Another practical point is temporal scope. Sequential recommendation combines long-range tendency and short-range fluctuation, so a single routing granularity can be suboptimal. RouteRec addresses this by applying the same MoE interface at multiple stages with different evidence windows. This keeps notation light while allowing routing behavior to adapt across horizons.

In the final model used in our main experiments, this multi-stage principle is realized with a shared router composition across all stages, while stage differences are introduced through temporal evidence windows and routing granularity. The next section formalizes this design.

4 Method

4.1 RouteRec Overview

Figure 2 summarizes RouteRec. We retain a strong sequential backbone and replace shared feed-forward computation with stage-wise MoE routing. The core interface follows the preliminary setup: hidden states are transformed by experts, and routing is controlled by behavior cues.

Let $\mathcal{T} = \{\text{macro}, \text{mid}, \text{micro}\}$ be the stage set. At stage $s \in \mathcal{T}$, routing uses a stage-specific behavior feature $\mathbf{f}_t^{(s)}$, while experts consume the shared hidden state $\mathbf{h}_t$:

$$\boldsymbol{\pi}_t^{(s)} = \operatorname{softmax}\Big(\operatorname{Router}_s(\mathbf{f}_t^{(s)})/\tau_s\Big), \tag{2}$$

$$\tilde{\mathbf{h}}_t^{(s)} = \sum_{e=1}^{E} \pi_{t,e}^{(s)} \operatorname{Expert}_e^{(s)}(\mathbf{h}_t). \tag{3}$$

Eq. (3) is the common stage interface. In the final RouteRec variant, all stages instantiate this interface with the same routing composition, while temporal differences arise from stage windows and routing granularity.

4.2 Behavioral Features from Raw Signals

RouteRec is designed for realistic data availability. It uses signals that are commonly present in interaction logs: item identity, timestamp, optional category, and optional corpus-level popularity. We intentionally avoid heavy side-information encoders in the routing branch. Instead, raw signals are converted into compact scalar summaries that are easy to compute, stable across datasets, and interpretable at analysis time.

For session s_m and target position t, the stage windows are

$$\mathcal{R}_t^{\text{macro}} = [s_{\max(1,m-H)}, \ldots, s_{m-1}], \tag{4}$$

$$\mathcal{R}_t^{\text{mid}} = [x_1, \ldots, x_{t-1}], \tag{5}$$

$$\mathcal{R}_t^{\text{micro}} = [x_{\max(1,t-w)}, \ldots, x_{t-1}], \tag{6}$$

where $H = 5$ and $w = 5$ in our main setting. A shared summary operator $\mathcal{A}(\cdot)$ maps each window to a low-dimensional feature vector:

$$\mathbf{f}_t^{(s)} = \mathcal{A}(\mathcal{R}_t^{(s)}), \qquad s \in \mathcal{T}.$$

The summary dimensions are grouped into four semantic families: *Tempo*, *Focus*, *Memory*, and *Exposure*. Tempo captures pace and interval statistics. Focus captures concentration and switching behavior over coarse categories or item groups. Memory captures recurrence, repetition, and overlap with recent context. Exposure captures interactions with popularity strata and their local drift. This decomposition is semantic rather than architectural: it gives routing modules an interpretable behavioral basis.

Importantly, the feature template is shared across stages, but each stage reads it through a different temporal lens. Macro features reflect persistent tendency and broad preference shape. Mid features track evolving intent inside the current session. Micro features emphasize immediate cues around the prediction target. This shared-template, stage-specific-view design keeps the feature branch compact while preserving temporal specificity.

4.3 Unified Router Composition

The final RouteRec architecture uses one shared scalar-group conditional composition at every stage. Experts are partitioned into G groups with C experts per group, so $E = GC$.

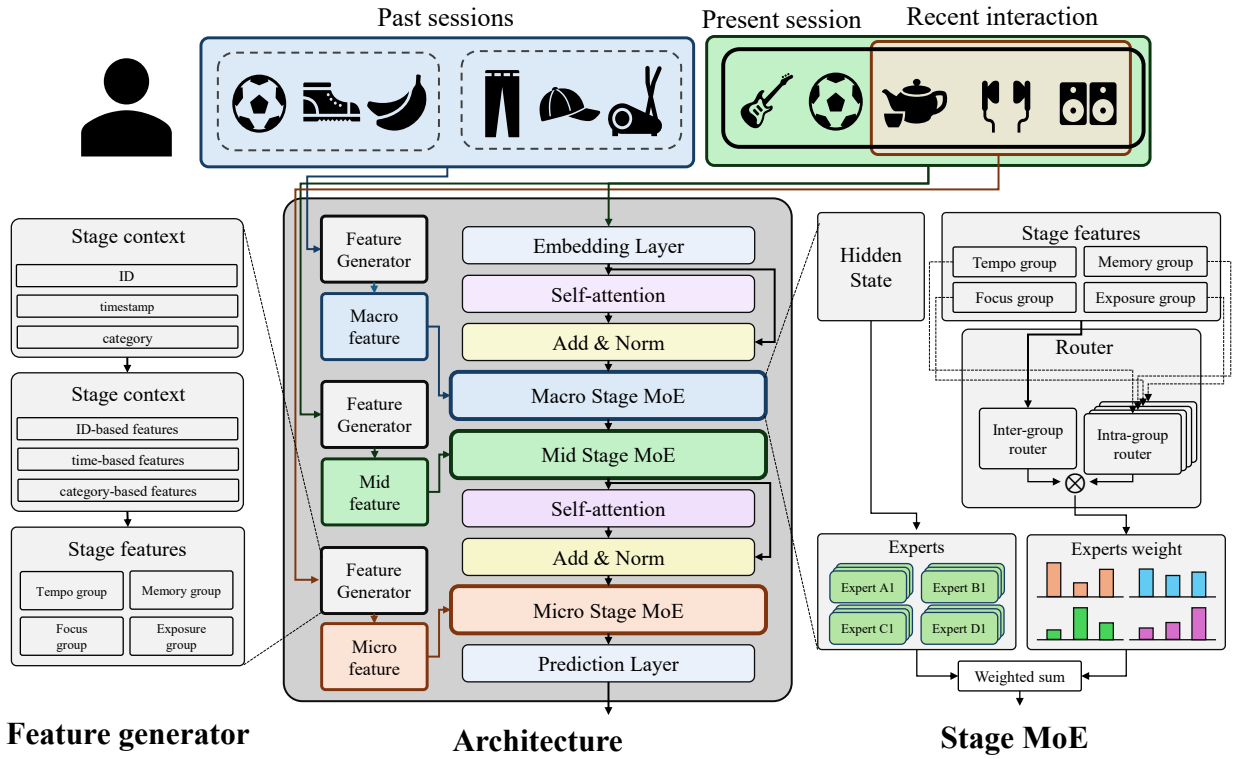

Figure 2: Overall RouteRec framework. (a) RouteRec keeps a strong sequential backbone and replaces shared feed-forward blocks with behavior-guided MoE stages. (b) All stages use the same scalar-group conditional routing composition: a group-level distribution and a conditional intra-group distribution are multiplied to form expert probabilities. (c) Macro, mid, and micro differ by temporal evidence scope and routing granularity, not by router formula.

For stage s and position t, the router produces

$$p_{t,g}^{(s)} = \mathrm{softmax}_g\left(z_{t,g}^{(s)}\right), \tag{7}$$

$$r_{t,g,c}^{(s)} = \mathrm{softmax}_c\left(u_{t,g,c}^{(s)}\right), \tag{8}$$

where $p_{t,g}^{(s)}$ is a group-level routing distribution and $r_{t,g,c}^{(s)}$ is a conditional within-group routing distribution. The final expert probability is their product:

$$\pi_{t,g,c}^{(s)} = p_{t,g}^{(s)}\, r_{t,g,c}^{(s)}, \tag{9}$$

followed by flattening (g, c) to expert index e. Eq. (9) is the same for macro, mid, and micro.

This unified composition is important for two reasons. First, it gives all stages a shared structural prior, which reduces stage-specific router engineering and simplifies debugging. Second, it keeps diagnostics comparable across stages because differences in routing behavior can be attributed to evidence scope rather than formula changes.

In RouteRec, routing is driven by behavior cues, while hidden states remain dedicated to expert transformation in Eq. (3). This preserves the routing-representation decoupling principle introduced in Preliminary.

4.4 Stage Instantiation by Temporal Scope

Although the router formula is shared, stages are not redundant. They differ in temporal evidence and routing granularity.

Macro and mid (session-level routing). Macro and mid compute one routing distribution per sequence from their respective windows, then broadcast it over valid positions inside the sequence. Macro uses broader cross-session history; mid uses within-session prefix evidence. This setting emphasizes stable allocation for slower-changing context.

Micro (token-level routing). Micro computes routing per token near the prediction target using short-horizon context. This enables fast adaptation to immediate transitions without changing the router form itself.

Formally, if Γ_s denotes the granularity operator at stage s, then

$$\boldsymbol{\pi}_t^{(s)} = \Gamma_s\left(\mathcal{R}_t^{(s)}\right),$$
$$\Gamma_{\text{macro}}, \Gamma_{\text{mid}} : \text{session-level},$$
$$\Gamma_{\text{micro}} : \text{token-level}.$$

Hence, RouteRec's coarse-to-fine behavior comes from stage evidence scope and update frequency, while the underlying router composition remains constant.

4.5 Training Objective

RouteRec is trained with a standard next-item objective and two router-side regularizers: route consistency and z-loss. The first regularizer shapes routing semantics, and the second stabilizes routing numerics.

Let $\mathbf{z}$ be final logits and y be the target item. The main loss is

$$\mathcal{L}_{\mathrm{CE}} = -\log \frac{\exp z_y}{\sum_j \exp z_j}.$$

For route consistency, we first compute a sample-level routing distribution at each stage by averaging token-level routing weights:

$$\bar{\boldsymbol{\pi}}_i^{\,s} = \frac{1}{L_i} \sum_{t=1}^{L_i} \boldsymbol{\pi}_{i,t}^{\,s}$$

where L_i is the valid sequence length. For each stage s, we construct a pair set $\mathcal{P}_s$ from feature-similar samples in the mini-batch and minimize

$$\mathcal{L}_{\mathrm{cons}} = \frac{1}{|\mathcal{T}|} \sum_{s \in \mathcal{T}} \frac{1}{|\mathcal{P}_s|} \sum_{(i,j) \in \mathcal{P}_s} \mathrm{JSD}\big(\bar{\boldsymbol{\pi}}_i^{\,s}, \bar{\boldsymbol{\pi}}_j^{\,s}\big)\,,$$

which encourages neighboring behavioral states to produce compatible routing distributions. In other words, semantically similar inputs should not induce arbitrarily different routing behavior.

To stabilize gating numerics, we adopt z-loss on router logits [38]:

$$\mathcal{L}_z = \frac{1}{|\mathcal{T}|} \sum_{s \in \mathcal{T}} \frac{1}{N_s} \sum_{(i,t) \in \mathcal{V}_s} \left(\log \sum_{e=1}^{E} \exp q_{i,t,e}^{\,s} \right)^2,$$

where $\mathcal{V}_s$ is the set of valid positions and $N_s = |\mathcal{V}_s|$. This term discourages excessive logit scale and helps prevent unstable, overly sharp routing. The final objective is

$$\mathcal{L} = \mathcal{L}_{\mathrm{CE}} + \lambda_{\mathrm{cons}}\, \mathcal{L}_{\mathrm{cons}} + \lambda_z\, \mathcal{L}_z,$$

where λ_{cons} and λ_z are trade-off coefficients. This objective jointly optimizes prediction quality, semantic consistency of routing, and numerical stability of expert selection.

5 Experiments

5.1 Experimental Setup

We evaluate RouteRec on six public datasets. These datasets cover e-commerce, location recommendation, short-video consumption, music listening, movie preference, and anonymous clickstream behavior. The six corpora are Beauty, Foursquare, KuaiRec, LastFM, ML-1M, and Retail Rocket. Appendix A summarizes their provenance and processed statistics.

Across all datasets, interactions are first ordered chronologically and then partitioned into sessions using dataset-specific inactivity thresholds. We adopt a 30-minute threshold for Foursquare, KuaiRec, LastFM, ML-1M, and Retail Rocket, whereas Beauty uses a 14-day threshold because review events are substantially sparser than click, play, or check-in logs. After sessionization, we perform iterative 5-core-style filtering until convergence: sessions of length below 5 are removed, items with fewer than 3 interactions are discarded, and session lengths are re-evaluated after item pruning. Beauty is constructed from Amazon Beauty reviews and additionally excludes interactions with rating below 3.

For transparency, the reported KuaiRec and LastFM results are obtained from processed subsets rather than the full raw corpora: specifically, we use a 20% KuaiRec subset and a 3% LastFM subset. Table 2 therefore reports statistics for the final processed datasets used in the experiments. After preprocessing, sessions are sorted by start time and divided at the session level into 70%/15%/15% train/validation/test partitions. In the initial split, the earliest 70% of sessions are assigned to training, and the remaining tail is divided into validation and test using stratification over target-repeat status, target-popularity bucket, and session-length bucket. For the final evaluation files, the training split is kept fixed, the held-out sessions are merged and re-ordered by session start time, and the merged tail is then divided contiguously into equal validation and test halves. During this held-out cleanup step, rows whose items are unseen in training are removed unless the unseen item is the final target of the session; unseen target cases are retained. When category or time signals are incomplete, RouteRec falls back to the same portable Tempo, Focus, Memory, and Exposure feature families computed from the available interaction log.

The baselines include GRU4Rec, SASRec, TiSASRec, DuoRec, FEARec, BSARec, FAME, DIF-SR, and FDSA [4, 8, 11, 14, 15, 21, 24, 33, 37]. All methods are trained within the same optimization family, with matched early-stopping rules and comparable capacity ranges. In the comparison tables, each baseline is represented by the run attaining the highest mean validation score over the nine reported metrics, and we report the corresponding test metrics. RouteRec is selected from the `fmoe_n4` family according to the same criterion.

5.2 Overall Performance: Does Behavior-Guided Routing Help?

Table 1 presents the main experimental comparison in the text using HR@10, NDCG@10, and MRR@20. This presentation keeps the main paper focused on a compact set of representative cutoffs, while Appendix A reports the full $k \in \{5, 10, 20\}$ grid. RouteRec exhibits its clearest gains on Beauty and the KuaiRec split, remains competitive on Retail Rocket, and stays close to the strongest shared-backbone baselines on Foursquare, LastFM, and ML-1M.

Two points are worth emphasizing. First, the improvement pattern is strongest when the target slice is behaviorally concentrated, which matches the motivation for behavior-guided routing. Second, the gain does not reduce to simple capacity scaling: several baselines already use rich sequence encoders or feature channels, yet RouteRec still benefits when the conditional computation path is chosen from explicit behavior cues rather than shared uniformly.

Dataset	Metric	SASRec	GRU4Rec	TiSASRec	FEARec	DuoRec	BSARec	FAME	DIF-SR	FDSA	RouteRec	Improv.
Beauty	HR@10	0.1347	0.0716	0.1318	0.1261	0.1146	0.0516	0.0315	0.1060	0.1347	**0.1616**	+20.0%
	NDCG@10	0.0748	0.0340	0.0878	0.0696	0.0680	0.0316	0.0211	0.0685	0.0909	**0.0982**	+8.1%
	MRR@20	0.0588	0.0247	0.0770	0.0544	0.0569	0.0260	0.0190	0.0592	0.0805	**0.0822**	+2.0%
Retail Rocket	HR@10	0.5605	0.5211	0.5435	0.5441	0.5377	0.4703	0.4697	0.4365	0.5133	**0.5612**	+0.1%
	NDCG@10	0.4099	0.3858	0.3963	0.3935	0.3864	0.3963	0.3934	0.3900	0.3651	**0.4150**	+1.2%
	MRR@20	0.3673	0.3479	0.3547	0.3515	0.3441	0.3743	0.3710	**0.3744**	0.3239	0.3737	-0.2%
Foursquare	HR@10	**0.3233**	0.2130	0.3157	0.3042	0.3197	0.2462	0.2546	0.3034	0.2998	0.3225	-0.2%
	NDCG@10	0.2014	0.1455	0.2016	0.1925	0.1955	0.1639	0.1646	0.1984	0.2001	**0.2041**	+1.3%
	MRR@20	0.1679	0.1278	0.1700	0.1619	0.1607	0.1410	0.1401	0.1690	**0.1723**	0.1712	-0.6%
ML-1M	HR@10	0.1478	0.1413	0.1538	0.1487	0.1520	0.1589	0.1557	0.1501	**0.1594**	0.1561	-2.0%
	NDCG@10	0.0770	0.0742	0.0844	0.0724	0.0811	**0.0854**	0.0853	0.0764	0.0835	0.0771	-9.7%
	MRR@20	0.0608	0.0587	**0.0690**	0.0550	0.0649	0.0678	0.0687	0.0595	0.0658	0.0585	-15.3%
LastFM	HR@10	**0.3994**	0.3090	0.3780	0.3765	0.3831	0.3453	0.3312	0.3617	0.3657	0.3794	-5.0%
	NDCG@10	**0.3267**	0.2673	0.3196	0.3167	0.3083	0.3024	0.2900	0.3149	0.3171	0.3253	-0.5%
	MRR@20	0.3054	0.2561	0.3027	0.2989	0.2863	0.2902	0.2784	0.3020	0.3033	**0.3099**	+1.5%
KuaiRec	HR@10	0.1055	0.0945	0.1011	0.1005	0.0995	0.0918	0.1011	0.0918	0.0967	**0.1060**	+0.5%
	NDCG@10	0.0943	0.0790	0.0916	0.0919	0.0919	0.0890	0.0915	0.0871	0.0865	**0.0965**	+2.3%
	MRR@20	0.0922	0.0753	0.0894	0.0903	0.0905	0.0882	0.0899	0.0863	0.0843	**0.0948**	+2.7%

Table 1: Main experimental results. For each baseline, we report the test metrics of the run achieving the highest mean validation score over the nine reported metrics. RouteRec is reported from the best run in the `fmoe_n4` family under the same selection criterion. Bold and underlined values denote the best and one second-best result within each row.

5.3 Router Input Study: Hidden vs. Feature vs. Both

We first isolate the router input itself. Figure 3(a) compares four controls: a shared FFN, hidden-only routing, feature-only routing, and mixed hidden+feature routing. Hidden-only routing improves over a shared FFN, but feature-only routing is stronger and more stable across datasets. Adding hidden states back into the feature router (both) does not consistently improve over feature-only routing.

This comparison supports a concrete claim: the gain is not from adding one more router branch, but from choosing a routing signal that reflects observable behavioral evidence.

5.4 Semantic Alignment: Does Feature Similarity Transfer to Routing Similarity?

Figure 3(b) analyzes whether nearby cases in feature space are also nearby in routing space. We bucket sample pairs by feature similarity and measure route similarity inside each bucket for hidden-only, both, and feature-only routing. Feature-only routing yields the clearest monotonic pattern: higher feature similarity corresponds to higher routing similarity.

This is the key semantic result behind the performance gap. Feature-guided routing is not only competitive in ranking metrics, but also better aligned with interpretable behavioral neighborhoods, which is exactly the intended control interface in RouteRec.

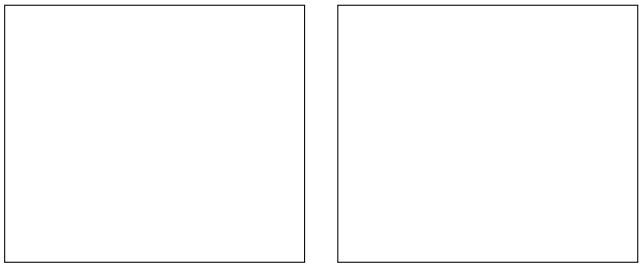

(a) Avg. ranking quality: FFN vs hidden vs both vs feature (b) Feature-similarity buckets vs route-similarity

Figure 3: Router-input and alignment analysis. (a) Average MRR@20 across shared FFN, hidden-only, mixed (hidden+feature), and feature-only routing. (b) Route similarity as a function of feature-similarity buckets, comparing hidden-only, mixed, and feature-only routing.

5.5 Is Coarse-to-Fine Staged Routing Necessary?

Figure 4 examines the RouteRec architecture itself. Panel (a) removes macro, mid, or micro routing to isolate the role of each temporal scope. Panel (b) separates staged routing from generic FFN replacement by comparing dense replacement against one-, two-, and three-stage layouts. Panel (c) tests wrapper-map and stage-order variants to show that the final composition is not an arbitrary design choice.

Across all three views, the same pattern is expected. Performance rises when routing is distributed across temporal scopes, dense replacement alone does not recover the same gain, and the final wrapper map remains among the strongest choices. The benefit of RouteRec therefore comes from coarse-to-fine behavioral allocation rather than from expert insertion by itself.

5.6 How Much Metadata Is Actually Needed?

Figure 5 evaluates whether RouteRec depends on rich metadata or whether ordinary interaction logs already provide enough routing evidence. The left panel removes cue families from the full model, and the right panel normalizes each variant by the gain of the full-cue model. Removing category-derived or time-derived cues lowers performance, but the sequence-only variant still retains a substantial fraction of the improvement over the shared-backbone baseline.

This result strengthens the portability claim. RouteRec benefits from richer evidence when it is available, but it does not require a heavy side-information stack. A meaningful part of the routing gain can already be recovered from the signals that most sequential recommendation systems already log.

Detailed full-cutoff results and extended analyses, including slice-wise and transfer studies, are reported in Appendices A–F.

6 Conclusion

RouteRec starts from a simple hypothesis: when behavioral evidence changes, the computation path should change as well. The experiments support this view from three main angles. First, RouteRec improves next-item ranking over strong sequential baselines across most datasets and cutoffs. Second, the gain comes from the routing signal and the staged design rather than from inserting experts indiscriminately. Third, lightweight cues derived from ordinary interaction logs already recover a substantial share of the benefit, which keeps the method practical even when side information is limited.

Taken together, these results suggest that compact behavioral cues provide a useful control interface for conditional computation in recommendation. Future work can build on this interface with broader transfer regimes, richer but still robust cue construction, and tighter deployment-aware constraints such as latency, continual adaptation, and router calibration.

References

[1] Shuqing Bian, Xingyu Pan, Wayne Xin Zhao, Jinpeng Wang, Chao Wang, and Ji-Rong Wen. 2023. Multi-modal Mixture of Experts Representation Learning for Sequential Recommendation. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management (CIKM)*. 110–119. doi:10.1145/3583780.3614978

[2] Oscar Celma. 2010. *Music Recommendation and Discovery in the Long Tail*. Springer. doi:10.1007/978-3-642-13287-2

[3] Yukuo Cen, Jianwei Zhang, Xu Zou, Chang Zhou, Hongxia Yang, and Jie Tang. 2020. Controllable Multi-Interest Framework for Recommendation. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*. 2942–2951. doi:10.1145/3394486.3403344

[4] Xinyu Du, Huanhuan Yuan, Pengpeng Zhao, Jianfeng Qu, Fuzhen Zhuang, Guanfeng Liu, Yanchi Liu, and Victor S. Sheng. 2023. Frequency Enhanced Hybrid Attention Network for Sequential Recommendation. In *Proceedings of the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*. 78–88. doi:10.1145/3539618.3591689

[5] William Fedus, Barret Zoph, and Noam Shazeer. 2022. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. *Journal of Machine Learning Research* 23, 120 (2022), 1–39. https://jmlr.org/papers/v23/21-0998.html

[6] Chongming Gao, Haonan Li, Yuchong Zhang, Meng Wang, Fuli Feng, Xiangnan He, Yong Li, and Depeng Jin. 2022. KuaiRec: A Fully-Observed Dataset and Insights for Evaluating Recommender Systems. *Proceedings of the 31st ACM International Conference on Information and Knowledge Management (CIKM)* (2022), 540–550. doi:10.1145/3511808.3557220

[7] F. Maxwell Harper and Joseph A. Konstan. 2015. The MovieLens Datasets: History and Context. *ACM Transactions on Interactive Intelligent Systems* 5, 4 (2015), 1–19. doi:10.1145/2827872

[8] Balázs Hidasi, Alexandros Karatzoglou, Linas Baltrunas, and Domonkos Tikk. 2016. Session-based Recommendations with Recurrent Neural Networks. In *International Conference on Learning Representations (ICLR)*. https://arxiv.org/abs/1511.06939

[9] Robert A. Jacobs, Michael I. Jordan, Steven J. Nowlan, and Geoffrey E. Hinton. 1991. Adaptive Mixtures of Local Experts. *Neural Computation* 3, 1 (1991), 79–87. doi:10.1162/neco.1991.3.1.79

[10] Michael I. Jordan and Robert A. Jacobs. 1994. Hierarchical Mixtures of Experts and the EM Algorithm. *Neural Computation* 6, 2 (1994), 181–214. doi:10.1162/neco.1994.6.2.181

[11] Wang-Cheng Kang and Julian J. McAuley. 2018. Self-Attentive Sequential Recommendation. In *2018 IEEE International Conference on Data Mining (ICDM)*. 197–206. doi:10.1109/ICDM.2018.00035

[12] Chao Li, Zhiyuan Liu, Mengmeng Wu, Yuchi Xu, Huan Zhao, Pipei Huang, Guoliang Kang, Qiwei Chen, Wei Li, and Dik Lun Lee. 2019. Multi-Interest Network with Dynamic Routing for Recommendation at Tmall. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management (CIKM)*. 2615–2623. doi:10.1145/3357384.3357814

[13] Jing Li, Pengjie Ren, Zhumin Chen, Zhaochun Ren, Tao Lian, and Jun Ma. 2017. Neural Attentive Session-based Recommendation. In *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management (CIKM)*. 1419–1428. doi:10.1145/3132847.3132926

[14] Jiacheng Li, Yujie Wang, and Julian J. McAuley. 2020. Time Interval Aware Self-Attention for Sequential Recommendation. In *Proceedings of the 13th ACM International Conference on Web Search and Data Mining (WSDM)*. 322–330. doi:10.1145/3336191.3371786

[15] Mingrui Liu, Sixiao Zhang, and Cheng Long. 2025. Facet-Aware Multi-Head Mixture-of-Experts Model for Sequential Recommendation. In *Proceedings of the 18th ACM International Conference on Web Search and Data Mining (WSDM)*. 127–135. doi:10.1145/3701551.3703552

[16] Qiao Liu, Yifu Zeng, Refuoe Mokhosi, and Haibin Zhang. 2018. STAMP: Short-Term Attention/Memory Priority Model for Session-Based Recommendation. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*. 1831–1839. doi:10.1145/3219819.3219950

[17] Ziwei Liu, Qidong Liu, Yejing Wang, Wanyu Wang, Pengyue Jia, Maolin Wang, Zitao Liu, Yi Chang, and Xiangyu Zhao. 2025. SIGMA: Selective Gated Mamba for Sequential Recommendation. *Proceedings of the AAAI Conference on Artificial Intelligence* 39, 12 (2025), 12264–12272. doi:10.1609/aaai.v39i12.33336

[18] Chen Ma, Peng Kang, and Xue Liu. 2019. Hierarchical Gating Networks for Sequential Recommendation. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*. 825–833. doi:10.1145/3292500.3330984

[19] Jiaqi Ma, Zhe Zhao, Xinyang Yi, Jilin Chen, Lichan Hong, and Ed H. Chi. 2018. Modeling Task Relationships in Multi-Task Learning with Multi-Gate Mixture-of-Experts. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge*

(a) Stage removal (b) Dense replacement vs. staged routing (c) Wrapper-map and order variants

Figure 4: Stage and layout ablations. (a) Stage removal. (b) Dense replacement versus staged routing. (c) Wrapper-map and stage-order variants.

(a) Avg. MRR@20 by cue family removal (b) Relative retention of the full-cue gain

Figure 5: Cue-family ablations. (a) Average MRR@20 under cue removal. (b) Gain retention relative to the full cue set.

Discovery and Data Mining (KDD). 1930–1939. doi:10.1145/3219819.3220007

[20] Julian McAuley, Rahul Pandey, and Jure Leskovec. 2015. Inferring Networks of Substitutable and Complementary Products. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*. 785–794. doi:10.1145/2783258.2783381

[21] Ruihong Qiu, Zi Huang, Hongzhi Yin, and Zijian Wang. 2022. Contrastive Learning for Representation Degeneration Problem in Sequential Recommendation. In *Proceedings of the 15th ACM International Conference on Web Search and Data Mining (WSDM)*. 813–823. doi:10.1145/3488560.3498433

[22] Retail Rocket. 2015. Retail Rocket Recommender System Dataset. https://www.kaggle.com/datasets/retailrocket/ecommerce-dataset

[23] Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc V. Le, Geoffrey E. Hinton, and Jeff Dean. 2017. Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer. In *International Conference on Learning Representations (ICLR)*. https://openreview.net/forum?id=B1ckMDqlg

[24] Yehjin Shin, Jeongwhan Choi, Hyowon Wi, and Noseong Park. 2024. An Attentive Inductive Bias for Sequential Recommendation beyond the Self-Attention. *Proceedings of the AAAI Conference on Artificial Intelligence* 38, 8 (2024), 8984–8992. doi:10.1609/aaai.v38i8.28747

[25] Fei Sun, Jun Liu, Jian Wu, Changhua Pei, Xiao Lin, Wenwu Ou, and Peng Jiang. 2019. BERT4Rec: Sequential Recommendation with Bidirectional Encoder Representations from Transformer. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management (CIKM)*. 1441–1450. doi:10.1145/3357384.3357895

[26] Hongyan Tang, Junning Liu, Ming Zhao, Xudong Gong, Zhaohui Jin, Yu Lei, Min Zhang, and Xuan Li. 2020. Progressive Layered Extraction (PLE): A Novel Multi-Task Learning (MTL) Model for Personalized Recommendations. In *Proceedings of the 14th ACM Conference on Recommender Systems (RecSys)*. 269–278. doi:10.1145/3383313.3412236

[27] Jiaxi Tang and Ke Wang. 2018. Personalized Top-N Sequential Recommendation via Convolutional Sequence Embedding. In *Proceedings of the 11th ACM International Conference on Web Search and Data Mining (WSDM)*. 565–573. doi:10.1145/3159652.3159656

[28] Chenyang Wang, Weizhi Ma, Chong Chen, Min Zhang, Yiqun Liu, and Shaoping Ma. 2023. Sequential Recommendation with Multiple Contrast Signals. *ACM Transactions on Information Systems* 41, 1 (2023), 1–27. doi:10.1145/3522673

[29] Shoujin Wang, Longbing Cao, Yan Wang, Quan Z. Sheng, Mehmet A. Orgun, and Defu Lian. 2021. A Survey on Session-Based Recommender Systems. *Comput. Surveys* 54, 7 (2021), 1–38. doi:10.1145/3465401

[30] Shoujin Wang, Liang Hu, Yan Wang, Longbing Cao, Quan Z. Sheng, and Mehmet A. Orgun. 2019. Sequential Recommender Systems: Challenges, Progress and Prospects. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence (IJCAI)*. 6332–6338. doi:10.24963/ijcai.2019/883

[31] Shu Wu, Yuyuan Tang, Yanqiao Zhu, Liang Wang, Xing Xie, and Tieniu Tan. 2019. Session-Based Recommendation with Graph Neural Networks. *Proceedings of the AAAI Conference on Artificial Intelligence* 33, 1 (2019), 346–353. doi:10.1609/aaai.v33i01.3301346

[32] Xu Xie, Fei Sun, Zhaoyang Liu, Shiwen Wu, Jinyang Gao, Bolin Ding, and Bin Cui. 2020. Contrastive Learning for Sequential Recommendation. arXiv:2010.14395 [cs.IR] doi:10.48550/arXiv.2010.14395

[33] Yueqi Xie, Peilin Zhou, and Sunghun Kim. 2022. Decoupled Side Information Fusion for Sequential Recommendation. In *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*. 1611–1621. doi:10.1145/3477495.3531963

[34] Dingqi Yang, Daqing Zhang, Zhiyong Yu, and Zhu Wang. 2013. A Sentiment-Enhanced Personalized Location Recommendation System. https://dl.acm.org/doi/10.1145/2507157.2507230 Foursquare check-in dataset used in many next-POI recommendation studies.

[35] Haochao Ying, Fuzhen Zhuang, Fuzheng Zhang, Yanchi Liu, Guandong Xu, Xing Xie, Hui Xiong, and Jian Wu. 2018. Sequential Recommender System based on Hierarchical Attention Networks. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence (IJCAI)*. 3926–3932. doi:10.24963/ijcai.2018/546

[36] Shengzhe Zhang, Liyi Chen, Dazhong Shen, Chao Wang, and Hui Xiong. 2025. Hierarchical Time-Aware Mixture of Experts for Multi-Modal Sequential Recommendation. In *Proceedings of the ACM Web Conference 2025 (WWW)*. 3672–3682. doi:10.1145/3696410.3714676

[37] Tingting Zhang, Pengpeng Zhao, Yanchi Liu, Victor S. Sheng, Jiajie Xu, Deqing Wang, Guanfeng Liu, and Xiaofang Zhou. 2019. Feature-level Deeper Self-Attention Network for Sequential Recommendation. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence (IJCAI)*. 4320–4326. doi:10.24963/ijcai.2019/600

[38] Barret Zoph, Alexander V. Kolesnikov, Neal Wadhwa, Jonathan Huang, Andrew M. Dai, Noam Shazeer, Quoc V. Le, and Yonghui Wu. 2022. Designing Effective Sparse Expert Models. arXiv:2202.08906 [cs.LG] doi:10.48550/arXiv.2202.08906

Dataset	Domain	Interactions	Sessions	Items	Avg. len.
Beauty	E-com.	33,488	4,243	3,625	7.9
Foursquare	POI	145,238	25,369	30,588	5.7
KuaiRec	Video	287,411	24,458	6,477	11.8
LastFM	Music	470,408	25,089	52,510	18.8
ML-1M	Movie	575,281	14,539	3,533	39.6
Retail Rocket	Browse	821,243	153,092	90,211	5.4

Table 2: Datasets used in the experiments.

A Dataset, Evaluation, and Full Results

This appendix summarizes the six datasets and the common evaluation protocol. The original releases are cited here for completeness [2, 6, 7, 20, 22, 34]. The dataset descriptions emphasize domain differences and signal availability because the routing branch depends on which behavioral evidence is exposed in each corpus.

A.1 Datasets and Protocol

Table 2 reports the statistics of the final processed datasets used in the experiments rather than the raw corpus sizes. In particular, KuaiRec corresponds to the 20% processed subset used in our runs, and LastFM corresponds to the 3% processed subset.

Beauty and Retail Rocket stress sparse e-commerce revisit behavior, Foursquare and KuaiRec emphasize stronger short-horizon intent changes, and LastFM and ML-1M provide longer preference-driven histories. This spread is useful because it tests whether behavior-guided routing helps only in highly dynamic settings or remains competitive when preference structure is more stable.

Signal availability also differs in a way that matters for routing. Beauty exposes item, time, and product-category cues after low-rating removal. Foursquare includes item, timestamp, and venue-category structure with bursty check-in dynamics. KuaiRec offers item, time, and coarse content-side metadata in short-video sessions with strong short-term intent shifts. LastFM provides repeated listening histories with artist-side signals, while ML-1M contributes longer preference-driven sequences with genre information. Retail Rocket is the sparsest side-information setting, relying mainly on anonymous event streams with item and time signals.

The preprocessing pipeline is unified across datasets in overall structure, while the session boundaries are adapted to the logging regime of each corpus. Foursquare, KuaiRec, LastFM, ML-1M, and Retail Rocket are sessionized with a 30-minute inactivity threshold, whereas Beauty uses a 14-day threshold because review timestamps are substantially more diffuse. After sessionization, we repeatedly remove sessions shorter than 5 and items with fewer than 3 interactions until no further pruning is required. This iterative filtering is necessary because item removal can render previously valid sessions too short, making a single-pass procedure insufficient.

Session splits are constructed at the session level after sorting by session start time. The earliest 70% of sessions are assigned to training, and the remaining 30% are treated as a held-out tail. This tail is first divided into validation and test using stratification over three factors: whether the target item already appears in the prefix, the target item's popularity bucket with respect to the training split, and the session-length bucket. We then build the final evaluation files by merging the held-out sessions, re-ordering them by session start time, and re-splitting them contiguously into equal validation and test halves. In these held-out files, non-target rows whose items are unseen in training are removed, whereas unseen target items are preserved. We report HR, NDCG, and MRR at $k \in \{5, 10, 20\}$ for all methods.

For the result tables, each baseline is represented by the run with the highest mean validation score over the nine reported metrics, and RouteRec is reported from the best run in the `fmoe_n4` family under the same selection rule.

A.2 Full Ranking Results

Table 3 restores the full cutoff grid summarized more compactly in Table 1. This presentation preserves the same selection rule as the main text while exposing the complete $k \in \{5, 10, 20\}$ breakdown.

The appendix table makes it easier to verify whether improvements survive beyond the compact main-text summary without changing the selection rule between the main text and appendix.

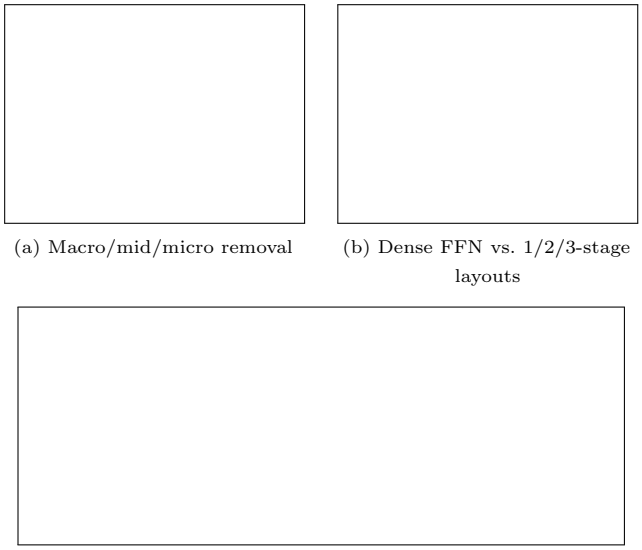

(a) Macro/mid/micro removal (b) Dense FFN vs. 1/2/3-stage layouts

(c) Wrapper-map and stage-order variants

Figure 6: Extended stage and layout ablations. (a) Stage removal. (b) Dense replacement versus staged routing. (c) Wrapper-map and stage-order variants.

B Extended Stage and Layout Ablations

Figure 6 expands the structural study in Figure 4 with the broader sweep of stage-removal and wrapper variants.

The figure should be read from left to right. Panel (a) verifies that macro, mid, and micro routing each contribute a distinct part of the gain. Panel (b) separates staged routing

Dataset	Metric	SASRec	GRU4Rec	TiSASRec	FEARec	DuoRec	BSARec	FAME	DIF-SR	FDSA	RouteRec	Improv.
Beauty	HR@5	0.0974	0.0372	0.0974	0.0974	0.0974	0.0430	0.0258	0.0802	0.0946	**0.1054**	+8.2%
	HR@10	0.1347	0.0716	0.1318	0.1261	0.1146	0.0516	0.0315	0.1060	0.1347	**0.1616**	+20.0%
	HR@20	0.1719	0.1032	0.1719	0.1719	0.1719	0.0602	0.0516	0.1433	0.1719	**0.2014**	+17.2%
	NDCG@5	0.0629	0.0227	0.0765	0.0603	0.0625	0.0289	0.0192	0.0602	0.0779	**0.0803**	+3.1%
	NDCG@10	0.0748	0.0340	0.0878	0.0696	0.0680	0.0316	0.0211	0.0685	0.0909	**0.0982**	+8.1%
	NDCG@20	0.0842	0.0418	0.0980	0.0810	0.0827	0.0339	0.0260	0.0778	0.1003	**0.1083**	+8.0%
	MRR@5	0.0514	0.0179	0.0694	0.0476	0.0506	0.0242	0.0170	0.0533	**0.0725**	0.0721	-0.6%
	MRR@10	0.0563	0.0226	0.0742	0.0515	0.0528	0.0253	0.0178	0.0567	0.0779	**0.0794**	+1.9%
	MRR@20	0.0588	0.0247	0.0770	0.0544	0.0569	0.0260	0.0190	0.0592	0.0805	**0.0822**	+2.0%
Retail Rocket	HR@5	0.4869	0.4475	0.4700	0.4647	0.4611	0.4413	0.4365	0.4305	0.4344	**0.4883**	+0.3%
	HR@10	0.5605	0.5211	0.5435	0.5441	0.5377	0.4703	0.4697	0.4365	0.5133	**0.5612**	+0.1%
	HR@20	**0.6316**	0.5863	0.6135	0.6200	0.6134	0.4986	0.5003	0.4409	0.5886	0.6313	-0.1%
	NDCG@5	0.3860	0.3620	0.3724	0.3678	0.3616	0.3869	0.3826	0.3881	0.3396	**0.3914**	+0.9%
	NDCG@10	0.4099	0.3858	0.3963	0.3935	0.3864	0.3963	0.3934	0.3900	0.3651	**0.4150**	+1.2%
	NDCG@20	0.4280	0.4023	0.4139	0.4127	0.4057	0.4034	0.4012	0.3911	0.3842	**0.4327**	+1.1%
	MRR@5	0.3524	0.3335	0.3399	0.3355	0.3284	0.3685	0.3644	**0.3733**	0.3082	0.3591	-3.8%
	MRR@10	0.3623	0.3434	0.3498	0.3462	0.3388	0.3724	0.3689	**0.3741**	0.3187	0.3688	-1.4%
	MRR@20	0.3673	0.3479	0.3547	0.3515	0.3441	0.3743	0.3710	**0.3744**	0.3239	0.3737	-0.2%
Foursquare	HR@5	0.2518	0.1743	0.2474	0.2418	0.2474	0.2014	0.2014	0.2426	0.2478	**0.2562**	+1.7%
	HR@10	**0.3233**	0.2130	0.3157	0.3042	0.3197	0.2462	0.2546	0.3034	0.2998	0.3225	-0.2%
	HR@20	**0.3833**	0.2590	0.3721	0.3645	0.3717	0.2866	0.3034	0.3521	0.3477	0.3777	-1.5%
	NDCG@5	0.1782	0.1331	0.1793	0.1721	0.1721	0.1496	0.1475	0.1786	**0.1834**	0.1827	-0.4%
	NDCG@10	0.2014	0.1455	0.2016	0.1925	0.1955	0.1639	0.1646	0.1984	0.2001	**0.2041**	+1.3%
	NDCG@20	0.2167	0.1572	0.2159	0.2077	0.2088	0.1741	0.1769	0.2107	0.2123	**0.2181**	+0.7%
	MRR@5	0.1540	0.1195	0.1568	0.1492	0.1473	0.1324	0.1297	0.1574	**0.1620**	0.1584	-2.2%
	MRR@10	0.1636	0.1246	0.1660	0.1578	0.1569	0.1382	0.1368	0.1656	**0.1689**	0.1673	-0.9%
	MRR@20	0.1679	0.1278	0.1700	0.1619	0.1607	0.1410	0.1401	0.1690	**0.1723**	0.1712	-0.6%
ML-1M	HR@5	0.0957	0.0906	0.1018	0.0883	0.0990	**0.1050**	0.0976	0.0888	0.1046	0.0939	-10.6%
	HR@10	0.1478	0.1413	0.1538	0.1487	0.1520	0.1589	0.1557	0.1501	**0.1594**	0.1561	-2.0%
	HR@20	0.2254	0.2119	0.2375	0.2277	0.2305	0.2282	0.2212	0.2300	**0.2430**	0.2337	-3.8%
	NDCG@5	0.0605	0.0581	0.0677	0.0531	0.0640	**0.0680**	0.0667	0.0566	0.0656	0.0572	-15.9%
	NDCG@10	0.0770	0.0742	0.0844	0.0724	0.0811	**0.0854**	0.0853	0.0764	0.0835	0.0771	-9.7%
	NDCG@20	0.0965	0.0919	**0.1054**	0.0924	0.1009	0.1028	0.1018	0.0963	0.1043	0.0964	-8.5%
	MRR@5	0.0489	0.0474	0.0565	0.0416	0.0525	0.0559	**0.0566**	0.0460	0.0529	0.0452	-20.1%
	MRR@10	0.0556	0.0539	0.0633	0.0495	0.0595	0.0631	**0.0642**	0.0542	0.0603	0.0534	-16.8%
	MRR@20	0.0608	0.0587	**0.0690**	0.0550	0.0649	0.0678	0.0687	0.0595	0.0658	0.0585	-15.3%
LastFM	HR@5	**0.3638**	0.2829	0.3522	0.3515	0.3526	0.3243	0.3097	0.3362	0.3430	0.3562	-2.1%
	HR@10	**0.3994**	0.3090	0.3780	0.3765	0.3831	0.3453	0.3312	0.3617	0.3657	0.3794	-5.0%
	HR@20	**0.4252**	0.3341	0.4038	0.3962	0.4089	0.3653	0.3508	0.3863	0.3909	0.4070	-4.3%
	NDCG@5	0.3152	0.2589	0.3112	0.3086	0.2984	0.2956	0.2832	0.3067	0.3097	**0.3178**	+0.8%
	NDCG@10	**0.3267**	0.2673	0.3196	0.3167	0.3083	0.3024	0.2900	0.3149	0.3171	0.3253	-0.5%
	NDCG@20	**0.3332**	0.2736	0.3260	0.3216	0.3149	0.3075	0.2950	0.3212	0.3234	0.3323	-0.3%
	MRR@5	0.2988	0.2509	0.2976	0.2942	0.2803	0.2860	0.2742	0.2969	0.2986	**0.3049**	+2.0%
	MRR@10	0.3037	0.2543	0.3011	0.2976	0.2844	0.2888	0.2770	0.3002	0.3016	**0.3079**	+1.4%
	MRR@20	0.3054	0.2561	0.3027	0.2989	0.2863	0.2902	0.2784	0.3020	0.3033	**0.3099**	+1.5%
KuaiRec	HR@5	0.0945	0.0842	0.0951	0.0951	0.0940	0.0902	0.0902	0.0891	0.0907	**0.1000**	+5.2%
	HR@10	0.1055	0.0945	0.1011	0.1005	0.0995	0.0918	0.1011	0.0918	0.0967	**0.1060**	+0.5%
	HR@20	**0.1235**	0.1109	0.1109	0.1158	0.1109	0.0923	0.1186	0.0995	0.1109	**0.1235**	+0.0%
	NDCG@5	0.0909	0.0756	0.0897	0.0902	0.0902	0.0885	0.0881	0.0863	0.0845	**0.0946**	+4.1%
	NDCG@10	0.0943	0.0790	0.0916	0.0919	0.0919	0.0890	0.0915	0.0871	0.0865	**0.0965**	+2.3%
	NDCG@20	0.0988	0.0831	0.0941	0.0958	0.0948	0.0892	0.0959	0.0891	0.0901	**0.1009**	+2.1%
	MRR@5	0.0897	0.0728	0.0879	0.0886	0.0890	0.0879	0.0874	0.0853	0.0824	**0.0928**	+3.5%
	MRR@10	0.0910	0.0742	0.0887	0.0893	0.0897	0.0882	0.0888	0.0857	0.0833	**0.0936**	+2.8%
	MRR@20	0.0922	0.0753	0.0894	0.0903	0.0905	0.0882	0.0899	0.0863	0.0843	**0.0948**	+2.7%

Table 3: Full experimental results across all cutoffs. Bold and underlined values denote the best and one second-best result within each row.

from generic expert insertion, and panel (c) checks that the final wrapper map is not an arbitrary configuration choice.

C Objective and Regularization Variants

Table 4 isolates the contribution of consistency, z-loss, and balance regularization using ranking quality, route consistency, and stability indicators.

This compact table shows whether the full objective improves MRR@20 while preserving routing consistency and stability.

D Extended Diagnostics

Figure 7 groups the main routing diagnostics into one coordinated layout so that stage usage, sharpness, and semantic consistency can be read together.

These diagnostics are organized from coarse structure to semantic interpretation. The first two panels show how broad or sharp routing is at each stage, while the last two panels connect those patterns back to feature similarity and interpretable behavioral regimes.

Table 4: Objective and regularization variants.

Variant	MRR@20	Consistency	Stability
No auxiliary loss	–	–	–
KNN + consistency	–	–	–
z-loss only	–	–	–
Balance loss only	–	–	–
Consistency + z-loss	–	–	–
Full objective	–	–	–

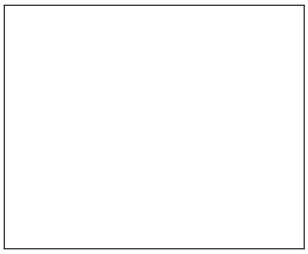

(a) Expert-usage heatmaps by stage

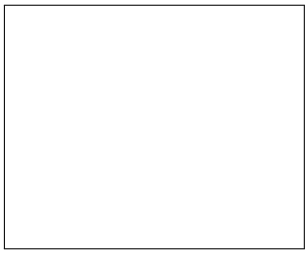

(b) Entropy and effective-expert-count summary

(c) Stage-wise route consistency

(d) Feature-bucket route patterns

Figure 7: Routing diagnostics. (a) Expert usage by stage. (b) Entropy and effective expert count. (c) Stage-wise route consistency. (d) Feature-bucket routing patterns.

E Behavior-Slice Analysis

Because RouteRec is motivated by heterogeneous behavior, a slice-wise view is a natural complement to the aggregate metrics. This appendix section is intended to localize gains to behavior regimes such as repeat-heavy prefixes, high-tempo sessions, narrow-focus sessions, and exploration-heavy prefixes.

The useful reading order is to first check which slices benefit most in panel (a), then verify in panel (b) whether those gains coincide with more decisive and semantically aligned routing. This experiment fits the paper's main argument especially well because it answers not only whether RouteRec helps, but also where the help appears.

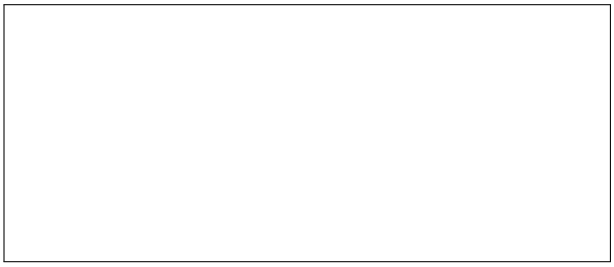

(a) Slice-wise MRR@20 across behavioral regimes

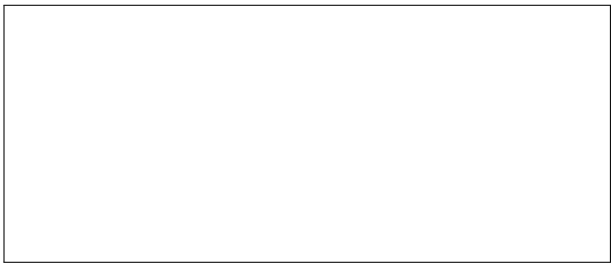

(b) Relative RouteRec gain and route concentration by slice

Figure 8: Behavior-slice analysis. (a) Slice-wise ranking quality. (b) Relative gain and routing concentration by slice.

Table 5: Source-target transfer matrix.

Src. → Tgt.	Bty	Fsq	KRec	LFM	ML1M	RR
Bty	–	–	–	–	–	–
Fsq	–	–	–	–	–	–
KRec	–	–	–	–	–	–
LFM	–	–	–	–	–	–
ML1M	–	–	–	–	–	–
RR	–	–	–	–	–	–

F Extended Transfer Results

This section reports additional transfer experiments for low-resource and cross-domain adaptation.

The matrix provides the exact source-target view that is too detailed for the main paper. It is useful for identifying asymmetric transfer pairs and for checking whether router reuse is most effective between domains with similar temporal and repetition structure. For compactness, the table uses the abbreviations Bty, Fsq, KRec, LFM, ML1M, and RR.

Together, the table and figure let the reader move from exact pairwise results to higher-level patterns. The intended appendix message is that transfer is asymmetric, most helpful in related domains and low-resource regimes, and partially preserved even under coarser router reuse.